

RETENTION INTELLIGENCE REVIEW

Churn & Retention Intelligence

Where the subscription book is losing customers and revenue, why those losses cluster where they do, and which accounts to defend first.

Analytical period Feb 2022 to Feb 2026 · snapshot reference date 1 March 2026

Book of 3,500 lifetime accounts · 2,547 active at snapshot · \$1,220,220 active monthly recurring revenue

Retention Analytics · Companion deliverable to the self-contained executive dashboard (executive-retention-command-center.html).

Findings are decision-support only, not audited financial results.

Contents

- Executive summary 4**
- Context and objectives 6**
 - The book under review 6
 - The decisions this report supports 6
- Data and methodology 7**
 - Data sources 7
 - How the synthetic data is built 7
 - Metric definitions 7
 - Governance and release gates 8
- Analytical framework 9**
 - Associative, not causal 9
 - The risk index 9
- Findings: retention health and trend 10**
 - The shape of the loss 11
- Findings: cohort retention 13**
- Findings: where churn concentrates 15**
 - By customer segment 15
 - By plan type 16
 - By acquisition channel 16
 - By region 17
- Findings: behavioural churn drivers 19**
 - Ranking the drivers by money, not just by rate 20
- Findings: revenue at risk and concentration 22**
 - Where the at-risk revenue sits 22
 - The risk tiers 23
- Risks, limitations, and caveats 25**
 - The data is synthetic 25
 - The recent months are right-censored 25

Revenue loss uses a monthly-value proxy	25
Relationships are associative	25
Tiering depends on the chosen thresholds	25
The book is observed at a single snapshot	25
Recommendations and action priorities	27
Sequencing, ownership, and success measures	28
Appendix	29
A. Cumulative churn share by commercial cut	29
B. Full behavioural driver detail	29
C. Recommended-action distribution across the scored book	30
D. Monthly retention trend, last twelve months	30
E. Figure index	30
F. Per-account feature dictionary	31

Executive summary

This review reads the full subscription history of a 3,500-account software book. Of those accounts, 953 have churned and 2,547 are still active at the 1 March 2026 snapshot, a cumulative customer loss of 27.2 percent of every account the business has ever signed. The active book carries \$1,220,220 of monthly recurring revenue. The question this report answers is not whether churn exists, which every subscription business lives with, but where it concentrates, what separates the accounts that leave from the accounts that stay, and how much current revenue sits behind each pattern.

27.2% CUMULATIVE CUSTOMER CHURN	\$1,220k ACTIVE MRR AT SNAPSHOT	223 CRITICAL + HIGH-RISK ACCOUNTS	\$191k WEIGHTED MRR EXPOSURE
---	---	---	--

Three findings carry the report. First, churn is not spread evenly. It concentrates sharply by who the customer is and how they were acquired. Startup-segment accounts have lost 44.8 percent of their number against 7.1 percent for Enterprise, a gap of more than six to one. Affiliate-sourced accounts churn at 42.1 percent against 14.9 percent for Partner-sourced accounts. Basic-plan accounts churn at 44.7 percent against 4.6 percent for Enterprise plans. The book does not have a churn problem so much as it has a low-end acquisition problem that turns into a churn problem.

Second, pre-churn account health separates leavers from stayers with unusual clarity. Accounts that carry a low net promoter score churn at 99 percent against 3 percent for accounts without that flag. Accounts with a heavy recent support burden churn at 94 percent. Failed payments, weak feature adoption, and declining usage each mark out populations that leave at several times the book rate. These signals are observed before the account closes, which is what makes them useful for prevention rather than only for post-mortem accounting.

Third, the revenue at stake is concentrated enough to act on without a broad campaign. The top 20 percent of churned accounts by value account for 57 percent of all lost monthly revenue. On the forward book, 24 accounts sit in the critical risk tier and 199 more in the high tier, together holding \$125,465 of monthly recurring revenue. Four intervention queues, ranked by weighted exposure, cover the material risk: a renewal save desk (\$123k weighted exposure across 898 accounts) leads, followed by payment rescue, adoption reactivation, and service recovery.




Figure 1. Retention plays ranked by weighted MRR exposure. The shaded bar is the monthly revenue in scope for each play; the solid bar is that revenue weighted by the historical churn rate of the accounts it targets.

The recommended sequence follows directly from these three findings and is set out in full in Section 11. In short: stand up the renewal save desk first because it carries the largest weighted exposure and the accounts are already identifiable; fix the payment-failure path second because it is cheap to run and the affected accounts churn at more than seven times the rate of accounts that pay cleanly; and rebalance acquisition spend away from the Affiliate and Paid Search channels third, because the cheapest lasting reduction in churn is to stop acquiring the accounts most likely to leave.

One caveat frames everything that follows. This analysis runs on deterministic synthetic data and the risk score is a transparent prioritisation index, not a calibrated churn probability. The patterns are internally consistent, but the exact magnitudes are illustrative. Section 10 states the limitations in full.

Context and objectives

A subscription business lives or dies on the difference between the revenue it adds and the revenue it loses. New logos and expansion get the attention, but retention sets the ceiling on everything else: a book that loses a quarter of its accounts over its life has to refill that quarter before it grows a single point. The purpose of this work is to give the people who own that problem, the revenue leadership and the customer success organisation, a clear and defensible picture of where the book is leaking and a ranked list of where to put their hands first.

The economics make the case on their own. Acquiring a replacement account costs a multiple of what it costs to keep an existing one, and the replacement starts at zero tenure with none of the expansion that a retained account accrues over time. A point of churn avoided is therefore worth more than a point of new logo growth at equal cost, because the retained account is already past its acquisition spend and its riskiest early months. This is the lens the report applies throughout: every avoidable departure is measured not only as lost monthly revenue but as acquisition spend the business will have to repeat to stand still.

The book under review

The dataset is the complete account history of a business-to-business software company: 3,500 accounts acquired across four customer segments (Enterprise, Mid-Market, SMB, and Startup), four regions (North America, Europe, APAC, and LATAM), six acquisition channels, and four plan tiers. Monthly activity runs from Feb 2022 to Feb 2026. At the snapshot date, 2,547 accounts are active and 953 have closed. The active book bills \$1,220,220 per month. Each account carries the behavioural signals a real customer success team would have in front of it: recent product usage and its trend, feature adoption, support ticket volume, net promoter score, and payment history.

The decisions this report supports

The work is built to answer five questions that a retention owner actually has to make calls on, rather than to produce a general description of the data:

- **How is the book trending?** Are monthly customer and revenue churn rising, flat, or falling, and is the loss weighted toward high-value or low-value accounts?
- **Where does churn concentrate?** Which segments, plans, channels, and regions carry a disproportionate share of the loss, and are recent cohorts retaining better or worse than older ones?
- **What separates leavers from stayers?** Which observable account-health signals are most strongly associated with churn, and by how much do they lift the rate?
- **Where is revenue exposed?** How much current MRR sits behind each pattern, and is the loss concentrated in a small set of high-value accounts or spread thinly across many small ones?
- **What should be done first?** Which intervention queues carry the largest weighted exposure, and in what order should they be resourced?

Everything in the findings maps back to one of these five questions. Where the data can only support an associative answer rather than a causal one, the report says so plainly rather than dressing a correlation as a cause.

Data and methodology

This section describes where the numbers come from and how each metric is defined, so that any figure in the report can be traced back to a rule rather than a judgement call. The pipeline is deterministic: it runs from a fixed seed, every artifact is reconstructible, and the analytical tables are governed by data contracts that are checked before any output is released.

Data sources

Four raw tables feed the analysis. A customer table holds the acquisition attributes (segment, region, channel, signup date). A subscription table holds the plan, the monthly recurring revenue, and the start and end dates that define whether and when an account churned. A weekly product-usage table records sessions and feature interactions. A payment table records billing events and failures. These are joined into one per-account feature row measured either at the account's churn date or, for active accounts, at the snapshot reference date, so that every account is observed at a comparable point in its own life rather than at an arbitrary calendar date.

Layer	Artifact	Grain	Role in this report
Raw	customers, subscriptions, usage, payments	event / account	Source of truth for all downstream tables
Features	customer_retention_features.csv	one row per account	Behavioural signals at churn or snapshot
Analysis	main_analysis_*, churn_by_*	aggregate	Trends, driver ranking, revenue at risk
Risk	customer_risk_scores.csv	active account	Tiered priority queue and recommended action
Cohort	cohort_retention_table.csv	cohort × month	Retention curves and heatmap

How the synthetic data is built

The data is generated, not collected, and it matters that the reader knows exactly what that means for the findings. The generator assigns each account a latent churn propensity driven by its segment, plan, channel, and a set of behavioural processes, then simulates usage, support, payment, and survey events consistent with that propensity. The result is a book in which the relationships between account health and churn are real and stable but designed rather than discovered. The report treats the magnitudes as illustrative; the metric definitions, the analytical framework, and the prioritisation logic are the transferable outputs.

Metric definitions

Several numbers in this report look similar and are easily confused, so each is pinned to an explicit definition here and used consistently throughout.

Metric	Definition
Cumulative customer churn	Closed accounts as a share of all accounts ever acquired. A lifetime, stock measure. Reads 27 percent here.
Monthly customer churn rate	Accounts that close in a month divided by accounts active at the start of that month. A flow measure, averaging near 1 percent over the window.
Cumulative churn share (by group)	Within a segment, plan, channel, or region, the share of that group's accounts that have closed.
MRR at risk	Current monthly recurring revenue billed by active accounts that carry the at-risk flag at the snapshot.
Churn rate lift	Churn rate inside a signal divided by churn rate outside it. A multiple, not a percentage.
Weighted MRR exposure	Revenue in an intervention queue multiplied by the historical churn rate of the accounts it targets. A heuristic for sizing, not a forecast.
Retention priority score	A transparent index combining behavioural risk and customer value, used only to rank the queue.

Two of these deserve emphasis. Cumulative customer churn and the monthly churn rate are different lenses on the same book and will not match: the first asks what share of everyone ever signed has left, the second asks how fast the active book is losing accounts right now. Both appear in this report and neither is wrong. Revenue loss throughout uses a monthly-value proxy, the account's recurring monthly revenue at the point it left, rather than a full contract or lifetime value, which the synthetic data does not model.

Governance and release gates

No figure in this report or the dashboard is released until the pipeline passes its governance gates. The artifacts are checked against data contracts that assert the shape, types, ranges, and referential integrity of every table, and against a final validation suite covering data quality, metric correctness, analytical integrity, the dashboard render, and release policy. At the current build all 46 of 46 checks pass, which is what licenses the report to make decision-support claims rather than screening-grade ones. The gate breakdown is below; the full check log lives in outputs/tables and is regenerated on every run.

Validation gate	Checks passing
Analytical Integrity	6 / 6
Dashboard Review	14 / 14
Data Quality	9 / 9
Governance & Release	7 / 7
Metric Correctness	10 / 10

The release policy is deliberately conservative about what these passes entitle the analysis to claim. Clearing every gate makes the output technically valid and analytically acceptable for decision support with explicit caveats. It does not make it committee-grade, because the synthetic-data limitation in Section 10 is an unresolved caveat by construction. The governance layer encodes that distinction rather than leaving it to the reader's goodwill.

Analytical framework

The analysis moves through four lenses in a deliberate order, each one narrowing the question. It begins with the trend, which tells you whether the book is getting better or worse in aggregate. It then decomposes the loss by commercial attribute and by acquisition cohort, which tells you where the loss lives. It then turns to account-level behaviour, which tells you what the leaving accounts had in common before they left. Finally it attaches revenue and a priority score to each pattern, which turns a diagnosis into a queue of accounts a team can work on Monday morning.

Associative, not causal

Every relationship in this report is associative. When the report says that accounts with a low net promoter score churn at a far higher rate, it means the two co-occur in the data, not that a low score causes the departure. The low score and the departure may both be downstream of a third thing, a poor onboarding or a product gap, and the data here cannot separate those. Establishing cause would require controlled experiments, holding out a population from an intervention and measuring the difference. The framework is built so that its outputs are the right inputs to exactly those experiments: the driver ranking tells you what to test, and the weighted exposure tells you what to test first.

The decomposition step is not cosmetic. A single book-wide churn rate can move in the opposite direction to every one of its parts if the mix of those parts is shifting, the trap known as Simpson's paradox. A business whose every segment is improving can post a worsening aggregate simply by acquiring more of its weakest segment. Cutting the rate by segment, plan, channel, and region guards against reading a mix shift as a performance change, and it is why the cohort and concentration findings are presented next to the aggregate trend rather than in place of it.

The risk index

Active accounts are scored on two axes and ranked on their product. The first axis is behavioural risk, built from the same pre-churn signals examined in Section 07: usage trend, feature adoption, support burden, net promoter score, and payment failures. The second axis is customer value, built from current monthly revenue and tenure. An account is worth working when it scores high on both, a high-value account showing real distress, rather than on either alone. The mean risk score across the 2,547 scored accounts is 12.7 on a 0-to-100 scale, which is deliberately low: most of the book is healthy, and the index exists to surface the minority that is not. Accounts fall into four tiers, critical, high, medium, and low, and each carries a single recommended action so the output is a worklist rather than a chart.

The index is transparent by design. Every input is a named, inspectable signal and the combination rule is a documented formula, not a black-box model. That transparency is a deliberate trade against predictive power: a trained classifier might rank marginally better, but a customer success manager can read this score, see why an account surfaced, and act on it without a data scientist in the room. For a prioritisation queue that property is worth more than a few points of accuracy.

Findings: retention health and trend

Across the analytical window the book lost an average of 0.7 percent of its active accounts per month and 0.3 percent of its active revenue per month. That revenue churn runs below customer churn is the single most reassuring number in the report: it means the accounts leaving are worth less, on average, than the accounts staying. The book is shedding its low end and keeping its high end, which is the favourable shape of churn to have. The later sections show exactly where that low end sits.




Figure 2. Monthly customer and revenue churn rates across the analytical window. The revenue line sits consistently below the customer line, the signature of low-value accounts leading the losses.

The level is healthy but the recent direction is not. Monthly customer churn in the final observed month reads 11.2 percent against a prior nine-month average of 1.1 percent. The last three months break clearly above the trend that preceded them. Two readings are possible and the report holds both. The operational reading is that something real has shifted in the most recent quarter and the book is losing accounts faster than it was. The measurement reading is that accounts closing near the snapshot date are mechanically over-represented because the window is right-censored, so the final months are not yet fully comparable to settled history. Section 10 returns to this. Either way the recent months warrant attention rather than alarm.




Figure 3. Monthly customer churn rate over the last twelve months. The final three months (highlighted) sit well above the prior nine-month average shown by the dotted line.

Underneath the rates, the active book has grown steadily in revenue terms to \$1,220,220 per month at the snapshot, so the business is adding more than it is losing. The retention question is therefore not about survival but about efficiency: every point of avoidable churn is a point the acquisition engine has to refill before the book grows, and the next sections show that a large share of the current loss is concentrated in places where it could be reduced.




Figure 4. Active monthly recurring revenue by month. The book has compounded steadily; the retention task is to protect the base, not to rescue a business in decline.

The shape of the loss

The gap between customer churn (0.7 percent a month on average) and revenue churn (0.3 percent) is worth reading deliberately, because its sign decides the whole posture of a retention programme. When revenue churn runs below customer churn, the average departing account is worth less than the average retained one, and the book is upgrading its quality even as it loses count. That is the case here, and it is the opposite of the dangerous pattern, where a business loses a few large accounts and posts low customer churn against high revenue churn. A retention owner who saw the dangerous pattern would drop everything to defend the top accounts; the owner here can instead run an efficient, mostly automated programme against a large low-value

tail and reserve human effort for the genuine high-value exceptions.

This shape also sets expectations for lifetime value. Because the accounts leaving are concentrated in the low-revenue, low-tenure end, the surviving book skews toward longer-lived, higher-value accounts, and blended lifetime value rises mechanically as the weak tail churns out. The risk in such a book is complacency: a healthy revenue-retention headline can mask a real and growing customer-count problem that will eventually starve the funnel. The cohort and concentration findings that follow are the check against that complacency.

SECTION 06

Findings: cohort retention

Aggregate churn rates hide whether the business is getting better at keeping the customers it acquires. Cohort analysis answers that by following each monthly intake of new accounts across its own life and comparing intakes at the same age. If recent cohorts retain better than older ones at month six, onboarding and product-market fit are improving. If they retain worse, the business is acquiring lower-quality accounts or onboarding them less well, and the aggregate rate will deteriorate later even if it looks stable now.




Figure 5. Retention by cohort age for the eight most recent monthly cohorts. Each line follows one acquisition month as its accounts age; the most recent cohorts are drawn most prominently.

Read at a fixed age, the curves bend the wrong way. Early cohorts hold almost all of their accounts through the first half-year; more recent cohorts start to separate downward at the same ages. The deterioration is modest in absolute terms, six-month retention still reads near 98 to 99 percent, but the direction is consistent and it is the direction that matters, because a cohort that retains one point worse at month six typically retains several points worse by the end of its second year. The honest qualifier is that the most recent cohorts are right-censored: they have not yet lived long enough to be compared fairly at the longer horizons, so the long-run gap should be read as a leading indicator to monitor rather than a settled fact.




Figure 6. Cohort retention heatmap. Each row is an acquisition month, each column a month of age, and each cell the share of the original cohort still active. Green is healthy retention, red is loss; the gradient down the early columns is where onboarding quality shows up.

The heatmap makes the pattern legible at a glance. The strongest retention sits in the upper rows, the older cohorts, and the early-age columns. As the eye moves down toward the recent cohorts, the early-age cells lose a little of their green. This is the same story the curves tell, shown as a surface rather than as lines, and it points the same way: the acquisition mix has been drifting toward accounts that are slightly harder to keep, which connects directly to the channel and segment findings in the next section.

One distinction matters for how this finding is used. Logo retention, the share of accounts still active, deteriorates a little faster than revenue retention, the share of original cohort revenue still billing, because the accounts leaving the recent cohorts are disproportionately the smaller ones. Six-month revenue retention holds near 99 percent even where logo retention slips below it. That is the same favourable shape seen in the aggregate trend, and it tells the retention owner that the cohort drift is a problem of small-account volume rather than an erosion of the cohorts' revenue spine. The fix belongs upstream in acquisition and onboarding, not in a high-touch save motion aimed at the recent cohorts.

SECTION 07

Findings: where churn concentrates

This is the section that turns a number into a target. The aggregate churn rate is an average over populations that behave nothing like each other, and averaging them together hides the only thing a retention team can act on: which populations are doing the leaving. Decomposed by segment, plan, channel, and region, the loss is strikingly uneven, and the unevenness lines up across the cuts in a way that points to a single underlying story.

By customer segment

Segment is the sharpest single cut. Startup accounts have lost 44.8 percent of their number and SMB accounts 34.0 percent, against 18.2 percent for Mid-Market and just 7.1 percent for Enterprise. The gradient is monotonic in account size: the smaller the customer, the more likely it is to leave. The same gradient runs through average revenue, where Enterprise accounts bill \$908 a month against \$170 for Startup, and through the health signals examined below. Small accounts churn more and are worth less individually, which is why the revenue line sits below the customer line in the trend.




Figure 7. Cumulative churn share by customer segment. The gradient is monotonic in account size, from Enterprise at the low end to Startup at the high end.

The segment story is not only about churn rates; it is about the whole health profile. The comparison below puts cumulative churn next to average net promoter score and the average usage trend for each segment. Enterprise accounts churn least, score highest on net promoter, and show usage that is still growing. Startup accounts invert all three. The signals are not independent symptoms; they move together, which is what makes a single segment-level health read meaningful.




Figure 8. Segment health on three axes: cumulative churn share, average net promoter score, and average usage trend. A segment that is losing accounts is also the segment with weak sentiment and declining usage.

By plan type

Plan tier tells the same story in pricing terms. Basic-plan accounts churn at 44.7 percent, nearly ten times the 4.6 percent rate of Enterprise-plan accounts, with Growth and Pro stepping down in between. Lower-priced plans attract lower-commitment buyers, and the data shows it. The commercial implication is not to abandon the entry tier, which feeds the funnel, but to recognise that a Basic-plan signup is a different and more fragile asset than a Pro-plan signup and should be onboarded and monitored as one.




Figure 9. Cumulative churn share by plan type. Churn falls steeply as plan tier rises.

By acquisition channel

Channel is where the analysis becomes directly actionable on spend. Affiliate-sourced accounts churn at 42.1 percent and Paid Search at 38.6 percent, against 14.9 percent for Partner and 18.8 percent for Referral. The accounts a business pays a third party to send it leave at roughly two to three times the rate of the accounts that arrive through a partner or a referral. This is the most controllable finding in the report, because acquisition mix is a budget decision. Shifting marginal spend from Affiliate and Paid Search toward Partner and Referral would lower the blended churn rate of every future cohort without touching a single retention playbook.




Figure 10. Cumulative churn share by acquisition channel. Paid and affiliate channels (red) deliver the least durable accounts; partner and referral channels (green) the most durable.

By region

Region is the weakest of the four cuts but not a flat one. LATAM accounts churn at 35.0 percent and APAC at 28.5 percent, both above the 27.2 percent book average, while North America sits below it at 23.8 percent. The spread is narrower than for segment or channel, and a good part of it probably reflects which segments and channels dominate each region rather than geography itself. Region is therefore a place to watch and to read alongside the other cuts, not a primary lever on its own.




Figure 11. Cumulative churn share by region against the book average. LATAM and APAC run hot; North America runs cool.

Taken together the four cuts tell one story from four angles. The accounts that leave are smaller, on cheaper plans, more often acquired through paid and affiliate channels, and somewhat concentrated outside North America. These are not four problems; they are four views of the same low-commitment population, and they explain both why revenue churn sits below customer churn and why the cohort mix has been drifting. The next section asks what those accounts looked like in the weeks before they left.

One word of caution carries over from the framework in Section 04. Because the four cuts overlap, they cannot simply be added. Affiliate traffic skews toward Startup and SMB accounts on Basic plans, so part of

the channel effect is really the segment and plan effect wearing a channel label, and part of the regional spread is which segments and channels dominate each region. Separating the independent contribution of each attribute would take a model that holds the others constant, which is the natural next step beyond this descriptive cut. For the decisions at hand the overlap does not change the conclusion: the levers point the same way whether the channel effect is direct or inherited, because shifting spend toward partner and referral also shifts the segment and plan mix toward the durable end.

SECTION 08

Findings: behavioural churn drivers

Commercial attributes say which accounts are structurally fragile. Behavioural signals say which accounts are in trouble right now, and they do it with the kind of separation that is rare in real data and worth examining carefully here. Each signal is a flag a customer success team can see in its own tools: a low recent net promoter score, a spike in support tickets, a failed payment, weak feature adoption, or a decline in usage.




Figure 12. Churn rate for accounts with and without each pre-churn signal. The gap between the two dots is the whole point: these flags split the book into populations that behave almost nothing alike.

The separation is stark. Accounts carrying a low net promoter score churn at 99 percent against 3 percent for accounts without it. Accounts with a heavy support burden churn at 94 percent against 4 percent. Failed payments mark a population that churns at 85 percent, weak feature adoption 69 percent, and declining usage 42 percent. Expressed as a multiple, a low net promoter score lifts the churn rate by more than 30 times and a heavy support burden by more than 26 times relative to accounts without the flag.

These multiples are large partly because the comparison group, accounts without any distress flag, almost never churns in this book. That is a property of synthetic data and the report does not lean on the exact magnitudes. What survives the caveat is the ranking and the practical consequence: sentiment and support signals identify near-certain churn, while usage decline identifies a much larger but lower-certainty population. A retention programme should treat the two differently, reserving expensive human save motions for the high-certainty sentiment and payment flags and using cheaper automated nudges for the broad usage-decline population.

Expressed as a multiple against accounts without each flag, the same five signals rank in a stable order. Low sentiment and high support burden sit far out at the top, payment failure in the middle, and weak adoption and usage decline closer to the no-lift line, though still well above it. The lift view and the absolute-separation view agree on the ranking and disagree only on emphasis, which is the reassurance that the ordering is a property of the data rather than of the chosen comparison.




Figure 13. Churn rate lift by behavioural signal, the ratio of the churn rate inside each signal to the rate outside it. Bars past the dotted no-lift line mark signals that raise the rate; the strongest sit several times above it.

Ranking the drivers by money, not just by rate

A high lift on a tiny population is less important than a moderate lift on a large and valuable one. The ranking below therefore scores each driver by its excess MRR association: the monthly revenue tied to churned accounts carrying the signal, above what the book-average churn rate would predict. This reorders the list. Support burden and low sentiment stay near the top because they combine a high rate with real revenue, but usage decline rises sharply, because although its lift is the smallest of the behavioural signals it touches the most accounts and the most money.




Figure 14. Churn drivers ranked by excess MRR association. Behavioural signals (red) dominate the structural attributes (slate); the label shows each driver's churn-rate multiple.

Driver	Accounts	Churn rate	Lift	Excess MRR assoc.
High Support Ticket	910	94%	3.5×	\$118,597
Low NPS	877	99%	3.6×	\$94,813
Usage Decline	1,842	42%	1.6×	\$92,386
Failed Payment	771	85%	3.1×	\$91,004
Low Feature Adoption	875	69%	2.6×	\$31,619

The five behavioural signals at the top of this table are the foundation of the risk index in Section 04 and of the intervention queues in Section 11. They are also the natural candidates for the controlled experiments the framework calls for: each one defines a population, a plausible intervention, and a revenue stake large enough to justify the test.

The signals also compound. An account that is both paying late and filing support tickets is in a different state from one carrying a single flag, and the risk index in Section 04 is built to capture that by scoring on the full signal set rather than on any one flag. In operational terms this means the save desk should not work five separate lists, one per signal, with the same account appearing on several. It should work one ranked list in which the multiply-flagged accounts rise to the top, and reserve its most expensive human contact for them. The single-flag, usage-decline-only accounts are the large base of the queue and belong in an automated motion, not a phone call.

SECTION 09

Findings: revenue at risk and concentration

Counting accounts tells you how many customers you are losing; it does not tell you how much money is leaving or whether it is worth a targeted defence. This section attaches revenue to the risk and asks the question that decides the operating model: is the exposure concentrated enough to defend account by account, or so diffuse that only a broad programme could touch it?

The answer is that it is highly concentrated. Ranking churned accounts by value, the top tenth carries 40 percent of all lost monthly revenue, the top fifth carries 57 percent, and the top third carries 69 percent. By the halfway point of churned accounts by value, 85 percent of all lost monthly revenue is already accounted for. A small set of departures does most of the financial damage, which is the single most important fact for designing the response: a focused save desk working a few hundred named accounts can address the majority of the revenue at stake.




Figure 15. Concentration of lost monthly value across churned accounts. The curve bows far above the line of even distribution; the marked point shows the top fifth of accounts carrying the majority of the loss.

Where the at-risk revenue sits

On the forward book, the same concentration appears across segments. The comparison below sets each segment's current at-risk MRR against the monthly value it has already lost to churn. SMB carries the largest current at-risk MRR at \$50,900, reflecting its size, while Enterprise, despite churning the fewest accounts, still carries \$36,630 of at-risk MRR because each Enterprise account is worth so much that even a handful in distress represents real money. This is why value has to sit alongside risk in the prioritisation: a single wavering Enterprise account can outweigh a dozen Startup accounts.




Figure 16. Current at-risk MRR against churned monthly value by segment. SMB leads on current exposure; Enterprise punches above its account count on value per account.

Segment	At-risk accts	At-risk MRR	Churned accts	Churned monthly value
Enterprise	42	\$36,630	38	\$24,186
Mid-Market	103	\$37,539	177	\$45,742
SMB	228	\$50,900	481	\$86,351
Startup	105	\$17,125	257	\$34,868

The risk tiers

The risk index sorts the 2,547 active accounts into four tiers. The critical tier holds 24 accounts billing \$11,464 a month, the high tier 199 accounts billing \$114,001, and the medium tier 468 accounts billing \$168,714. The critical and high tiers together, 223 accounts and \$125,465 of MRR, are small enough for a team to work by hand and large enough to matter. They are the spine of the priority queue.




Figure 17. Account count and MRR exposure by risk tier. The critical and high tiers are a workable number of accounts holding a material share of exposed revenue.

The score distribution behind those tiers confirms that the book is mostly healthy and the risk is a tail, not a bulk. The low tier dominates the count and clusters at the bottom of the score range; the critical and high tiers are a thin, well-separated tail at the top. That shape is exactly what a prioritisation index should produce: it is not labelling half the book as at risk, it is isolating the minority that genuinely is.




Figure 18. Distribution of churn risk scores by tier. The mass sits in the healthy low tier; the at-risk tiers form a distinct upper tail.

SECTION 10

Risks, limitations, and caveats

A report that hides its weaknesses is less useful than one that names them, because the reader cannot judge how far to trust a finding without knowing what could undermine it. Five limitations bound everything above and each changes how a specific finding should be read.

The data is synthetic

The book is generated from a fixed seed, not collected from a live business. The relationships are real and stable but designed, so the magnitudes are illustrative. The behavioural lifts in Section 08 in particular are larger than a real book would show, because the comparison group of fully healthy accounts almost never churns here. The method, the pipeline, the metric definitions, and the prioritisation logic carry over to real data without change; the specific percentages do not.

The recent months are right-censored

The apparent acceleration in the final three months of the trend, where the rate climbs to 11.2 percent, is partly a measurement artifact. Accounts that close near the snapshot are mechanically over-represented relative to accounts that will close later but have not yet, so the tail of the window overstates the settled rate. The finding to carry forward is that the recent months merit monitoring, not that the book has entered a churn crisis.

Revenue loss uses a monthly-value proxy

Every revenue figure is built from the account's recurring monthly revenue at the point it left or at the snapshot, not from contract value or modelled lifetime value. This understates the true cost of losing a long-contract Enterprise account and overstates the relative importance of easily-replaced monthly accounts. The concentration finding in Section 09 is robust to this, since both ends use the same proxy, but absolute exposure figures should be read as monthly run-rate, not annual contract value.

Relationships are associative

No finding in this report establishes cause. A low net promoter score does not necessarily cause a departure; it co-occurs with it. The intervention queues in the next section are therefore presented as hypotheses to be validated with controlled tests, not as guaranteed returns. The weighted exposure numbers size the prize; they do not promise it.

Tiering depends on the chosen thresholds

The value tiers and risk cut-points are quantile-based and would need recalibration on a real distribution. Moving a threshold moves accounts between tiers and changes the headline counts. The ranking is stable to these choices; the exact tier populations are not, and any operational commitment to a tier size should be set against real data.

The book is observed at a single snapshot

Every active account is measured at one date, the reference snapshot, rather than followed continuously. That keeps the cross-sectional comparison clean, because every account is read at a comparable point in its own

life, but it means the analysis cannot see seasonality, recovery, or the path an account took to its current state. An account that spiked on support tickets last month and has since settled looks the same at the snapshot as one still deteriorating. A live deployment would refresh the snapshot on a schedule and watch the trajectory between refreshes, which would sharpen the risk index and catch recoveries the single snapshot misses.

SECTION 11

Recommendations and action priorities

The findings converge on a short, ordered list. The ordering is by weighted exposure and by speed to value, not by how interesting the finding is, and each recommendation names the finding it rests on, the population it targets, and the metric that would prove it worked.

The four intervention queues that anchor the priorities are sized below. Read the weighted-exposure column as the order of magnitude of monthly revenue each play is defending, and the scope column as the gross revenue passing through the queue. The gap between the two reflects the historical churn rate of the targeted accounts: the renewal save desk works a large book at a moderate rate, while payment rescue works a small book at a very high rate, which is why the second is the stronger return on effort even though the first defends more revenue in absolute terms.

Intervention queue	Accounts	MRR in scope	Hist. churn	Weighted exposure
Renewal Save Desk	898	\$291,539	42%	\$122,965
Payment Rescue	119	\$50,156	85%	\$42,415
Adoption Reactivation	287	\$23,317	69%	\$16,098
Service Recovery	46	\$9,763	94%	\$9,162

P1 Stand up a renewal save desk for Mid-Market and SMB

This play carries the largest weighted exposure in the book at \$122,965, across 898 accounts holding \$291,539 of MRR in scope, with the focus on Mid-Market and SMB where the value and the volume meet (Section 09, Figure 1). Work the critical and high risk tiers first, lead with the accounts flagged on sentiment and support, and run a pre-renewal save playbook with tailored offers. Measure success as the saved MRR among contacted accounts against a held-out control, not as gross contact volume.

P1 Fix the payment-failure path

Accounts with a failed payment churn at 85 percent, more than 8 times the rate of accounts that pay cleanly (Section 08). The payment-rescue queue is only 119 accounts but carries \$42,415 of weighted exposure, and the fix is cheap: dunning sequencing, card-update prompts, and retry logic. This is the highest return-on-effort action in the report and should ship in parallel with the save desk. Measure recovered MRR per failed-payment account.

P2 Rebalance acquisition spend away from the weakest channels

Affiliate and Paid Search accounts churn at 42.1 and 38.6 percent against 14.9 percent for Partner and below 20 percent for Referral (Section 07, Figure 10). Shifting marginal budget toward partner and referral motions lowers the blended churn rate of every future cohort with no retention effort at all. Measure the six-month retention of new cohorts by channel and track the blended rate down.

P2 Run an adoption-reactivation programme for the usage-decline tail

Usage decline touches 1,842 accounts, the largest behavioural population, at a lower per-account certainty than the sentiment flags (Section 08). It suits a cheaper, automated motion: onboarding refreshers, feature nudges, and success check-ins. The reactivation queue carries \$16,098 of weighted exposure across 287 accounts. Measure the share of flagged accounts that return to a positive usage trend.

P3 Treat low-tier and Basic-plan signups as a distinct, instrumented class

Basic-plan accounts churn at 44.7 percent (Section 07). Rather than a save motion, this calls for a product and onboarding fix: a guided first-value path for entry-tier accounts and clearer upgrade prompts for the ones that show real engagement. Measure the rate at which Basic accounts either reach an activation milestone or upgrade within ninety days.

P3 Convert the top drivers into controlled experiments

Because every relationship here is associative (Section 10), the driver ranking should feed a small experiment programme: hold out a control for each of the top intervention queues and measure the lift directly. This turns the weighted-exposure estimates into known returns and protects the budget from chasing a correlation that does not respond to treatment.

Sequenced this way, the first two actions defend the revenue already in distress and ship within a quarter; the middle two reduce the inflow of fragile accounts and the size of the at-risk tail over the following two quarters; and the last two harden the programme so its returns are measured rather than assumed. None of the six depends on a tool the team does not already have, and all six trace to a quantified finding in this report.

Sequencing, ownership, and success measures

The plan resolves to the following sequence. Each action carries one owner, one horizon, and one measure that settles whether it worked, so the programme can be governed on outcomes rather than activity.

#	Action	Owner	Horizon	Success measure
P1	Renewal save desk	Customer Success	0–90 days	Saved MRR vs held-out control
P1	Payment-failure path	Billing / RevOps	0–60 days	Recovered MRR per failed-payment account
P2	Acquisition rebalance	Marketing	1–2 quarters	6-month retention of new cohorts by channel
P2	Adoption reactivation	Customer Success	1–2 quarters	Share of flagged accounts returning to positive usage
P3	Basic-plan onboarding fix	Product	2 quarters	90-day activation or upgrade rate
P3	Driver experiments	Analytics	Ongoing	Measured lift vs control per queue

Appendix

A. Cumulative churn share by commercial cut

Cut	Group	Accounts	Churned	Churn share	Avg MRR
Segment	Enterprise	536	38	7.1%	\$908
Segment	Mid-Market	974	177	18.2%	\$467
Segment	SMB	1,416	481	34.0%	\$246
Segment	Startup	574	257	44.8%	\$170
Plan	Enterprise	371	17	4.6%	\$1,660
Plan	Pro	966	144	14.9%	\$514
Plan	Growth	1,207	365	30.2%	\$180
Plan	Basic	956	427	44.7%	\$62
Channel	Affiliate	392	165	42.1%	\$249
Channel	Paid Search	823	318	38.6%	\$316
Channel	Outbound	634	156	24.6%	\$490
Channel	Organic	777	169	21.8%	\$389
Channel	Referral	378	71	18.8%	\$401
Channel	Partner	496	74	14.9%	\$537
Region	LATAM	555	194	35.0%	\$383
Region	APAC	488	139	28.5%	\$367
Region	Europe	1,090	294	27.0%	\$411
Region	North America	1,367	326	23.8%	\$402

B. Full behavioural driver detail

Signal	Accounts	Churn w/ signal	Churn w/o signal	Lift
Low NPS	877	99%	3%	30.2×
High Support Ticket	910	94%	4%	25.7×
Failed Payment	771	85%	11%	7.7×
Low Feature Adoption	875	69%	13%	5.3×
Usage Decline	1,842	42%	10%	4.0×

C. Recommended-action distribution across the scored book

Recommended action	Accounts	Share of scored book
monitor only	1,953	76.7%
renewal conversation	437	17.2%
customer success outreach	80	3.1%
product adoption campaign	64	2.5%
executive save motion	7	0.3%
billing intervention	6	0.2%

D. Monthly retention trend, last twelve months

The active book at the start of each month, the accounts and revenue lost during it, and the resulting churn rates. The customer and revenue churn columns are the data behind Figures 2 and 3; the right-censoring caveat in Section 10 applies to the final rows.

Month	Active accts	Active MRR	Churned	Cust. churn	Rev. churn
Mar 2025	2,614	\$1,038k	6	0.2%	0.1%
Apr 2025	2,685	\$1,068k	7	0.3%	0.1%
May 2025	2,756	\$1,090k	11	0.4%	0.2%
Jun 2025	2,818	\$1,128k	14	0.5%	0.3%
Jul 2025	2,867	\$1,152k	18	0.6%	0.4%
Aug 2025	2,936	\$1,185k	37	1.3%	0.9%
Sep 2025	2,964	\$1,202k	41	1.4%	0.8%
Oct 2025	2,984	\$1,217k	57	1.9%	0.8%
Nov 2025	2,992	\$1,234k	89	3.0%	1.2%
Dec 2025	2,979	\$1,252k	136	4.6%	2.1%
Jan 2026	2,917	\$1,243k	217	7.4%	3.3%
Feb 2026	2,771	\$1,220k	309	11.2%	5.3%

E. Figure index

Figures 1 through 18 are the static export of the chart pack in outputs/graphs/, regenerated deterministically from the same processed data and analytical tables that drive the live dashboard. The dashboard (executive-retention-command-center.html) carries the interactive versions of the trend, segment, driver, and queue views, with period, segment, and channel filters. This report and that dashboard are two renderings of one governed pipeline; neither contains a number the other cannot reproduce.

F. Per-account feature dictionary

The behavioural signals examined in Sections 07 and 08 and scored in the risk index are measured at the account's churn date or, for active accounts, at the snapshot. The principal fields are defined below.

Field	Definition
tenure_days	Days from acquisition to churn or snapshot.
current_mrr	Account's monthly recurring revenue at observation.
usage_trend	Change in 30-day session count; negative is decline.
recent_sessions_30d / 90d	Product sessions in the trailing window.
feature_adoption_score_recent	Breadth of feature use, 0 to 100.
support_tickets_30d / 90d	Support tickets filed in the window.
nps_score_recent	Most recent net promoter response.
failed_payments_90d	Count of failed billing attempts in 90 days.
payment_failure_flag	Any failed payment in the window.
renewal_near_flag	Renewal falls within the near-term horizon.
at_risk_flag	Account carries at least one distress signal.
churn_flag	Account has closed (the analysis target).